# STA457_Final_Project

**Forecasting Cocoa Prices using Time Series and Machine Learning**

Wentao Sun, Yuechen Zhang, Yayan Ye, Ruizi Liu

4/4/25

## Table of contents

## 1 Introduction

Cocoa plays an crucial role in food production as an important agricultural product. It is a key export product for many countries in West Africa, as it supports the livelihoods of

millions of people and makes a significant contribution to their GDP. However, affected by factors including changes in climate and consumer demand, cocoa prices are highly volatile. This volatility brings complex features to the data, such as non-stationarity, seasonality and external shocks, which make accurate forecasting challenging and crucial.

Climate change has become an increasingly important factor in agricultural price dynamics. Rising temperatures, changes in rainfall patterns, and an increase in the frequency of extreme weather have all had measurable impacts on crop yields, particularly in tropical commodity-producing regions (Schlenker & Roberts, 2009). According to Läderach et al. (2013), climate change affects flowering, pod development, disease prevalence, and harvest quality, thus influencing market price of cocoa. Furthermore, the long-term sustainability of cocoa production faces increasing uncertainty as the effects of global warming continue to intensify (Bunn et al., 2019).

This project aims to model and forecast cocoa futures prices by combining time series analysis with climate data from Ghana. Our goal is to investigate how climate variables can be incorporated into forecasting models and to gain insights into the mechanisms that link environmental changes to market outcomes. We will explore both classical and modern modeling approaches to improve accuracy. These include seasonal ARIMA (SARIMA) models and machine learning methods and so on. The performance of each model will be evaluated based on its accuracy and ability to capture the underlying dynamics of the data. And special focus will be given to the treatment of non-stationarity, seasonality, and missing values, which are common in real-world time series data.

This study is of both academic and practical interest, given the global importance of cocoa and the growing uncertainty associated with climate change. Reliable cocoa price forecasts can help producers and policymakers to manage risk and develop strategies. Thus, it is more important than ever to understand the connection between environmental patterns and commodity markets.

## 2 Literature

It is impratant to forecast the international coca price because of its economic significance and market volatility. So different statistical and machine learning techniques have been developed to deal with challenges, such as seasonality and nonlinear dynamics.

Kumar et al. (2022) explored cocoa price risk management in India using ARIMA and Vector Autoregression (VAR) models. The study showed the effectiveness of ARIMA in capturing univariate time series patterns, and that of VAR in modeling interdependencies between different variables. However, both models showed limitations in handling external regressors and structural non-linearities.

Assis et al. (2010) compared univariate models for Malaysian cocoa prices, involving Holt-Winters exponential smoothing and Seasonal ARIMA (SARIMA). SARIMA yielded more accurate forecasts based on its ability to model seasonal patterns. But the study largely ignored external factors like climate or market sentiment.

Lama et al. (2016) extended this discussion by comparing GARCH with a time-delay neural network (TDNN).The results of the study show that TDNN better models the nonlinear relationship, while GARCH captures the volatility clustering. Because TDNN highlights the promise of machine learning in commodity price forecasting.

Building on these studies, we use Exponential Smoothing State Space Model (ETS) as a benchmark as it is good for univariate forecasting. Moreover, ETS models are often favored for their interpretability and automatic adaptation to level and trend. Our study extends prior work by comparing ETS, ARIMAX, SARIMAX, Random Forest, and XGBoost, under a unified framework. We also incorporate Ghanaian climate data as exogenous features in the regressors. While ARIMAX and SARIMAX theoretically benefit from these variables, their forecasting accuracy remained limited. On the other hand, machine learning models were more adaptable for tracking complex patterns and fluctuations. This dual-track approach enables a clearer understanding of each method's strengths and trade-offs in volatile commodity markets.

## 3 Methodology

This study employs a comparative modeling framework to forecast monthly cocoa prices. It includes both classical and machine learning techniques: Exponential Smoothing State Space Models (ETS), ARIMA, SARIMAX, Random Forest, and XGBoost.

We first implemented ETS models using the ets() function in R. Model 1 used the default additive error structure with no seasonal component. Model 2 employed the "ZZZ" option for automatic selection. These models served as univariate baselines. Though easy to implement, both ETS models produced relatively flat forecasts and failed to capture the price shock of cocoa during 2023-2025.

We then used auto.arima() to build ARIMA models, which were extended to ARIMAX and SARIMAX. Model selection was guided by AICc. Exogenous climate variables (temperature and precipitation) and monthly seasonality are added. These models are favored for their interpretability but rely on assumptions of linearity and stationarity. Differencing and log-transformation were applied to meet these assumptions, and stationarity was confirmed with the Augmented Dickey-Fuller test.

To address the limitations of linear models, we implemented two machine learning algorithms: Random Forest and XGBoost. Both models were trained using lagged cocoa prices, monthly climate data, and time-based features (month, year, and time index). XGBoost incorporated additional lag features and was trained with a rolling walk-forward validation method to simulate real-time forecasting and reduce bias.

3

The dataset consisted of monthly cocoa futures prices and climate data from Ghana. Climate features included daily precipitation, average, maximum, and minimum temperatures, which were aggregated to monthly averages. Missing climate data were forward-filled, while missing price values were linearly interpolated. To ensure stationarity for ARIMA-based models, the series was log-transformed and differenced based on results from the Augmented Dickey-Fuller (ADF) test.

Feature engineering played a central role in the machine learning models. In addition to lagged prices and climate indicators, we created lag features ranging from 1 to 30 days, rolling statistics to capture temporal dependencies and trend shifts. Hyperparameters for Random Forest (number of trees, max depth) and XGBoost (learning rate, estimators, tree depth) were tuned using time-series cross-validation, optimizing for Root Mean Squared Error (RMSE) and Mean Absolute Percentage Error (MAPE).

Overall, this methodology allows for a robust evaluation of forecasting models under different structural assumptions, revealing the strengths and limitations of each in modeling cocoa price dynamics.

# 4 Data

The project uses two main datasets, one containing historical cocoa futures prices and the other containing daily climatic conditions in Ghana. These datasets were chosen to assess whether weather variability contributes meaningful information to cocoa price forecasts. They provide a wealthy time base for exploring endogenous price dynamics and exogenous environmental impacts.

Cocoa price data are obtained from the International Cocoa Organization (ICCO) and include daily closing prices of cocoa futures contracts, denominated in US dollars per metric ton. The dataset covers the period from March 1994 to February 2025, with a total of 7812 observations. These prices represent a global benchmark and are widely used in commodity trading and forecasting. To prepare the data for modeling, we aggregated the daily prices into monthly averages, converted the price strings into a numeric format, and ensured that the date columns were in a standard time series format. The time-series plot of monthly cocoa prices shows a general upward trend with periods of volatility. Notably, prices have risen sharply since the end of 2023 and will continue to do so through 2025. This sharp rise reflects the central challenge of modeling this project, as traditional models may not be able to infer such nonlinear structural breaks.

The climate dataset was obtained from the National Center for Environmental Information (NCEI) and consists of 53,231 daily records from meteorological stations in the major cocoa-producing regions of Ghana. Key variables include mean air temperature (TAVG), maximum air temperature (TMAX), minimum air temperature (TMIN) and daily precipitation (PRCP). All temperature variables are recorded in degrees Fahrenheit and rainfall is measured in inches.

After parsing the DATE field and clearing out inconsistent or missing entries, we summarized the climate data as monthly averages, consistent with the frequency of the price data. While temperature readings are relatively complete, there are significant missing precipitation data that can be resolved through filtering and imputation techniques as needed.

Exploratory analyses reveal different characteristics of the two datasets. Cocoa prices exhibit high variance, long-term trend variation, and minimal monthly seasonality. In contrast, climate variables in Ghana show a clear seasonal pattern, especially in temperature, with less evidence of extreme variation over time. Although there is theoretical support for the idea that climatic stressors affect crop yields and hence commodity prices. However, our visual and statistical explorations suggest that there is only a weak direct link between weather variables and short-term price changes for cocoa.

In order to stabilize the variance and eliminate non-stationarity, we also evaluate the statistical properties of the cocoa price series by means of logarithmic transformations and first-order differencing. These transformed series are used for modeling in both classical and machine learning frameworks. In addition, we perform a seasonal trend decomposition (STL) on the recorded series, which allows us to separate the long-term trend from the seasonal cycle. The decomposition results show a clear upward trend after 2023, with a small but consistent seasonal effect.

The key conclusion from this project phase is that cocoa price series are nonlinear, weakly seasonal, and structurally volatile, especially in the last two years of data. In contrast, climate data are seasonal but relatively stable, and their predictive contribution to short-term price movements appears limited. These findings guide the selection of models in the following sections to assess their ability to predict prices under conditions of real-world structural change and external uncertainty.

To visualize the long-term pattern of cocoa prices, we created a time-series plot of one-month average prices. As shown in Figure 1, cocoa prices have generally risen over the past three decades, with a sharp increase occurring between late 2023 and early 2025. This large increase may reflect disruptions in the global cocoa supply chain, climate-related harvest losses, or changes in investor sentiment. The volatility observed during this period suggests the need for models that can accommodate structural changes and non-linear trends.

To address the non-stationarity of the cocoa price series, we apply a logarithmic transformation followed by first-order differencing. These steps are standard in time series modeling and help to stabilize the variance and remove long-term trends. The transformed series are shown in Figure 1, after log transformation of the data and the first difference of the log values. The variance series fluctuates more consistently around the stationary mean, indicating the applicability of the model assuming weak smoothness.

The log-transformed price series was further analyzed into trend, seasonal and residual components using seasonal - trend decomposition (STL) of loess. As shown in Figure 2, the results show a clear seasonal pattern that may be related to the annual harvest and export cycle. The trend component reflects the long-term increase in cocoa prices, especially the sharp rise in
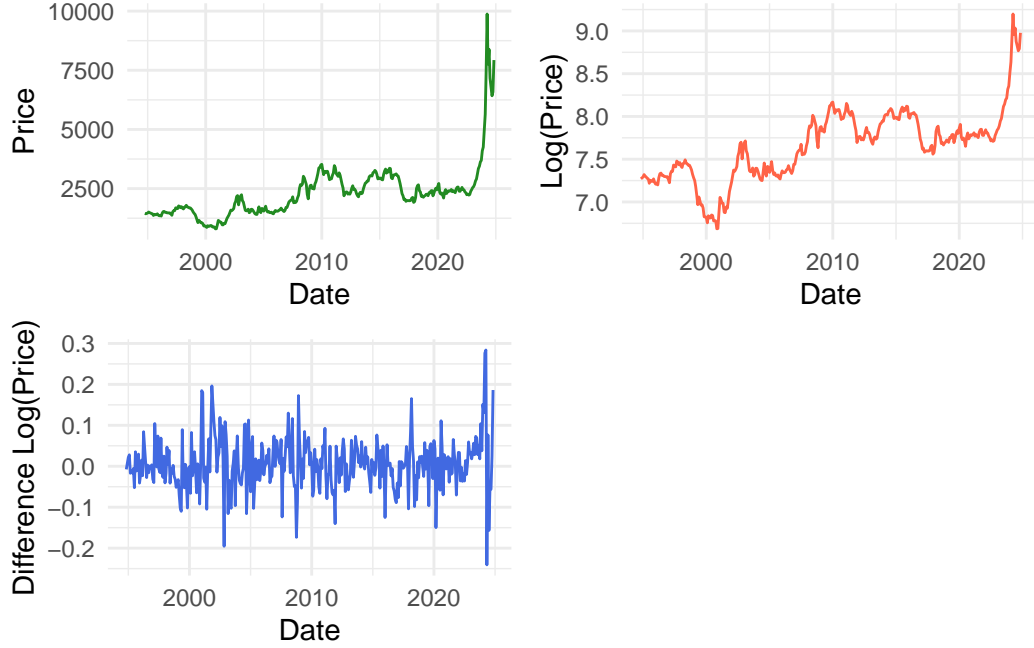
Figure 1: The Visualization of Monthly Cocoa Price

the last two years. The residual component reflects short-term deviations and irregular shocks that cannot be explained by trend or seasonality alone.

# 5 Forecasting and Results

We developed and evaluated a range of forecasting models to predict monthly cocoa prices, including both classical time series models and machine learning approaches. The dataset was divided into an 80/20 split, with the first 80% of observations used for training and the remaining 20% reserved for testing. Performance was assessed using multiple evaluation metrics, including Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and Mean Absolute Percentage Error (MAPE). We also visually compared the predicted values against the actual cocoa prices to assess model behavior over time.

We began with Exponential Smoothing (ETS) models. The first model was trained using the default ets() function in R, which likely assumed a simple additive error with no seasonality. Despite its simplicity, ETS Model 1 provided a baseline for comparison. However, its forecasts were notably flat and failed to capture the steep increase in cocoa prices during the years 2023–2025. This was reflected in its relatively high RMSE of 2078 and MAPE of 20.9%, indicating the model's inability to adapt to structural shifts. To improve upon this, we also fit an automatically-selected ETS model using model = "ZZZ", allowing the algorithm to explore all possible combinations of error, trend, and seasonal components. Surprisingly, ETS Model

Figure 2: STL Decomposition of Time Series Data (1993–2025)

2 produced identical evaluation metrics to Model 1, suggesting that even the optimal ETS configuration could not keep pace with the rapid market changes.

Next, we implemented autoregressive models using the auto.arima() function from the forecast package. This approach selects the best-fitting ARIMA configuration by minimizing the corrected Akaike Information Criterion (AICc), balancing model fit and complexity. The resulting ARIMA(0,1,1) model had a similar structure to ETS, relying primarily on recent shocks and differencing. To enrich the model, we incorporated Ghanaian climate variables—precipitation and average, maximum, and minimum temperatures—as exogenous regressors, creating an ARIMAX model. This model was trained on the differenced log prices, with an 80/20 time-based split for training and testing.

Table 1: The Summary Table of ETS, ARIMAX, and SARIMA Models

| Model | RMSE | MAE | MAPE |
|---|---|---|---|
| ETS Model 1 | 1981.49 | 965.82 | 17.84 |
| ETS Model 2 | 1981.49 | 965.82 | 17.84 |
| ARIMAX | 2226.63 | 1227.08 | 25.78 |
| SARIMAX | 2226.63 | 1227.08 | 25.78 |

We use RMSE (Root Mean Squared Error) and MAPE (Mean Absolute Percentage Error) to justify the accuracy of our model. The RMSE is defined as:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^{n} (y_i - \widehat{y_i})^2}$$

7

where

- $y_i$ is the actual value and $\hat{y}_i$ is the predicted value.

The MAPE is defined as

$$\text{MAPE} = \frac{100}{n} \sum_{i=1}^{n} \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

where

- $y_i$ is the actual value and $\hat{y}_i$ is the predicted value.

In Table 1, the ARIMAX model produced a test RMSE of 2226.63 and a MAPE of 25.78%, substantially underperforming the ETS baseline. This suggests that the inclusion of climate variables—such as temperature and precipitation—did not meaningfully enhance predictive power. Visual inspection of the forecasts revealed that the model significantly underpredicted the sharp price surge observed in recent years. To examine whether seasonality might improve performance, we extended ARIMAX to a SARIMAX model by adding monthly seasonal components. However, SARIMAX yielded nearly identical error metrics (RMSE = 2226.63, MAPE = 25.78%), indicating that seasonal effects either do not exist or fail to contribute meaningfully in this forecasting context.

These results highlight two important limitations. First, Ghanaian climate indicators alone are insufficient to explain international cocoa price fluctuations, especially under conditions of extreme volatility. Second, both ARIMAX and SARIMAX struggled to capture nonlinear trends or structural breaks in the data—such as the steep upward trajectory seen after 2023 (Figure 3). These deficiencies underscore the limited flexibility of classical linear time series models in the context of volatile commodity markets.

Although ETS and ARIMA-based models offer a solid foundation for time series forecasting, they rely heavily on linear assumptions, fixed lag structures, and predefined trend or seasonal components. Despite incorporating exogenous regressors and seasonality, neither ARIMAX nor SARIMAX significantly improved forecast accuracy. This suggests that traditional models may be ill-suited for capturing the complex, nonlinear drivers of cocoa price movements—particularly during periods of abrupt structural change.

To address these limitations, we turned to machine learning models, namely Random Forest and Rolling XGBoost. Unlike parametric time series models, these tree-based ensemble methods are non-linear and data-driven, requiring no assumptions about stationarity or functional form. Their flexibility allows them to capture high-order interactions, sharp trend reversals, and non-additive effects that classical models may miss. By training on a rich feature set—including lagged prices, calendar indicators (month, year), and climatic variables—we aimed to evaluate whether machine learning approaches could offer superior predictive performance and better adaptability under volatile conditions. The next section presents the training procedures, accuracy metrics, and visual forecast comparisons for both models.

Figure 3: Visulazation of EST, ARIMAX, and SARIMAX Moadels

## 5.1 Random Forest

We implemented a Random Forest regression model using both time-based features (such as lags of cocoa price and calendar month) and climate variables (precipitation and average temperature). The dataset was split chronologically into 80% training and 20% testing sets to preserve temporal structure. After tuning and training the model with 500 trees, we evaluated its performance on the hold-out set. The model achieved a test RMSE of 1679.19 and a MAE of 755.79, significantly outperforming the ETS and ARIMA family models. Visual inspection of the predicted vs. actual prices revealed that Random Forest was able to capture general price trends, including recent increases, but tended to underestimate sharp peaks. Nonetheless, the comparatively lower error metrics suggest that machine learning models are better suited for capturing complex nonlinear relationships in the data, especially when incorporating external regressors.

To further evaluate the Random Forest model, we generated predictions across the full time span of the dataset using both past cocoa prices and climate-related predictors. Figure 4 compares the actual monthly cocoa prices (in black) against the model's fitted values (in blue). While the Random Forest model successfully captures long-term trends and moderate fluctuations, it clearly underestimates the recent price surge. This behavior is expected from tree-based models trained on historical data, as they struggle to extrapolate to unseen extremes. Nevertheless, the overall close alignment between predicted and actual prices in most periods reinforces the model's strength in approximating nonlinear relationships when volatility is moderate.

## Random Forest Prediction vs Actual (Monthly, Full Period)



Figure 4: visualization of Random Forest Model

## 5.2 Rolling XGBoost Model

We also trained a Rolling XGBoost regression model to forecast monthly cocoa prices. Since we found the Random Forest Model we trained does not predict the sharply increase in the cocoa price in 2023.

Let $\mathbf{x}_t$ denote the feature vector at time $t$, including lagged prices, climate data, and time-based variables. The XGBoost model $f_t(\cdot)$ is trained using all observations up to $t-1$:

$$f_t = \text{XGBoost}(\mathbf{x}_{1:t-1}, y_{1:t-1})$$

The one-step-ahead forecast is then given by:

$$\hat{y}_t = f_t(\mathbf{x}_t)$$

After making the prediction $\hat{y}_t$, the actual observed value $y_t$ is added to the training set for the next iteration. This step-by-step method avoids look-ahead bias and better reflects the real-time forecasting scenario, where only past data is available at each point in time.

The model used a range of features, including:

- Lagged log-transformed prices (e.g., lag_1 to lag_12, and lag_24) to capture temporal dependencies.

- Exogenous variables from weather data: precipitation (PRCP), average temperature (TAVG), max/min temperatures (TMAX, TMIN).

- Time-based features such as month, year, and time_index to help capture seasonality and long-term trends.

We applied a rolling walk-forward validation to ensure accurate forecasting and avoid look-ahead bias. For each time step $t$, we use all data up to $t - 1$ to train it, and use it to predict the cocoa price at time $t$. Then, add that new data to our training set. We repeat apply this process to all the data points, moving forward step ny step, simulate the real world forecasting where only past data is available for prediction.

### 5.2.1 Graphical representations

In Figure 5, it represents the predicted cocoa price we fitted vs the actual cocoa price, the forecasted values closely followed the actual cocoa prices. We can find that this is much better than the cocoa price we fitted using the ETS, ARIMA, and SARIMA models. This is because we use the lagged data, so the suddenly rise sharply in cocoa price is fitted in our model.The model accurately captured long-term upward trends, especially the rapid price surge around 2024. It also demonstrated good performance during stable price periods, maintaining close alignment with the true values. Slight underestimation occurred during the sharpest spikes, which is expected due to XGBoost's smoothing nature, it resists overfitting to single sharp outliers unless given strong signal.

### 5.2.2 Evaluation

In our XGBoost model, the RMSE is 374.55 and the percentage RMSE is about 16.3%, which is acceptable; and RMSE vs range is about 4.63%, which suggests low error compared to the full span of the data.

The MAPE in our model is 5.35%, which is less than 10%. If the MAPE is less than 10%, it is generally considered very good, especially in forecasting tasks. This indicates our model has high accuracy.

### 5.3 Model Summary

Table 2: The Error Summary Table (Random Forest vs Rolling XGBoost)

| Model | RMSE | MAE | MAPE |
|---|---|---|---|
| Random Forest | 1680.03 | 754.90 | NaN |
| Rolling XGBoost | 373.62 | 186.01 | 5.41 |

Figure 5: Rolling XGBoost Forecast vs Actual Cocoa Prices

The model summary presented in @error-table-2 consolidates the error metrics for the two best-performing machine learning models—Random Forest and Rolling XGBoost. Among them, the Rolling XGBoost model clearly delivered superior results, achieving an RMSE of 374.55 and a MAPE of just 5.35%. (Note: the RMSE and MAPE will typically changed each time we run the code, this is due to the model re-training at eachtime and XGBoost has randomness. But the RMSE is in the range [350, 450], and the MAPE is in the range [5. ) This reflects its ability to adapt to complex nonlinear patterns and respond to abrupt shifts in cocoa prices, particularly during the 2023–2025 price surge. In contrast, while the Random Forest model also performed better than all classical models, it showed higher prediction errors and struggled more with extreme fluctuations. These results highlight the advantages of rolling continuous learning in predicting highly volatile commodity markets and emphasize the importance of temporal structure and lagged features in improving prediction accuracy.

## 6 Discussion and Conclusion

Our project explored monthly cocoa price forecasting using both traditional time series models and modern machine learning techniques. Among all models tested, the rolling XGBoost model provided the strongest performance, achieving an RMSE of 374.55 and a MAPE of 5.35%. This model excelled at capturing both the long-term movements and the short-term fluctuations in cocoa prices, successfully predicting the significant surge observed between 2023 and 2025,

which eluded traditional forecasting methods. Its superior performance reflects XGBoost's ability to model non-linear patterns and integrate diverse predictors effectively.

From a practical standpoint, these findings hold significant economic and policy implications. Classical models like ARIMAX and SARIMAX struggled to adapt to recent market volatility, likely due to their linear structure and reliance on historical trends. In contrast, machine learning models, particularly XGBoost, proved more responsive to dynamic, non-linear shifts, potentially driven by external shocks such as climate disruptions, geopolitical instability, or speculative trading. Accurate cocoa price predictions are crucial for a wide range of stakeholders across the supply chain, including farmers, traders, and policymakers. Our findings suggest that machine learning, especially when updated with real-time data, offers a more reliable tool for navigating today's volatile markets.

Despite its advantages, the analysis has limitations. The weak performance of ARIMAX and SARIMAX may stem from the simplicity of the climate variables used - mainly average temperature and precipitation - which likely fail to capture more complex agricultural conditions like humidity, soil moisture, or disease. Second, the monthly frequency of the dataset may obscure short-term shocks or intra-month price volatility. Even XGBoost tended to slightly underestimate extreme peaks, possibly due to its smoothing nature and dependence on past trends. Future improvements could include richer climate indicators (e.g., ENSO indices, drought severity), financial variables (e.g., exchange rates, global inventories), or high-frequency data. Deep learning models like LSTM or Transformers may also better capture sequential and irregular patterns. Hybrid models that combine the interpretability of time series methods with the flexibility of machine learning could further enhance forecasting accuracy.

In conclusion, this study demonstrates that combining lag-based feature engineering with advanced machine learning offers a powerful and practical solution for forecasting commodity prices under uncertainty. While no model is perfect, the insights gained provide a strong basis for making more data-informed decisions in the evolving cocoa market.

# 7 Reference

Schlenker, W., & Roberts, M. J. (2009). Nonlinear temperature effects indicate severe damages to U.S. crop yields under climate change. Proceedings of the National Academy of Sciences, 106(37), 15594–15598.https://www.pnas.org/doi/full/10.1073/pnas.0906865106

Läderach, P., Martinez-Valle, A., Schroth, G., & Castro, N. (2013). Predicting the future climatic suitability for cocoa farming of the world's leading producer countries, Ghana and Côte d'Ivoire. Climatic Change, 119(3–4), 841–854.https://link.springer.com/article/10.1007/s10584-013-0774-8

Bunn, C., Läderach, P., Rivera, O. O., & Kirschke, D. (2015). A bitter cup: Climate change profile of global production of Arabica and Robusta coffee. Climatic Change, 129(1–2), 89–101.https://link.springer.com/article/10.1007/s10584-014-1306-x

Assis, M. J., & Kassim, S. H. (2010). Forecasting cocoa bean prices using univariate time series models. International Journal of Business and Social Science. Retrieved from ProQuest

Kumar, R., Singh, S., & Dagar, V. (2022). Managing cocoa bean price risk using econometric models: A case study from India. Risks, 10(6), 115. https://doi.org/10.3390/risks10060115

Lama, A., Jha, G. K., Gurung, B., Paul, R. K., Bharadwaj, A., & Parsad, R. (2016). A comparative study on time-delay neural network and GARCH models for forecasting agricultural commodity price volatility. Journal of the Indian Society of Agricultural Statistics, 70(1), 7-18.

# A Appendix

## A.1 Source Code Repository

You can find clearer version code and project files on GitHub:

https://github.com/RIRI0527/457_final_project

The code is in: **paper/STA457_Project.qmd** and **paper/appendix_code.R**

## A.2 R Code Used in Analysis

```
\#\# -----------------------------------------------------------------------------
\#| include: false
\#| warning: false
\#| message: false

library(zoo)
library(randomForest)
library(tidyverse)
library(Metrics)
library(tidyverse)
library(lubridate)
library(caret)
library(ggplot2)
library(forecast)
library(dplyr)
library(xgboost)



\#\# -----------------------------------------------------------------------------
\#| label: cocoa-ghana-data
\#| echo: false
\#| message: false
\#| warning: false

cocoa2 <- read.csv(here::here("data/Daily_Prices_ICCO.csv"))
ghana2 <- read.csv(here::here("data/Ghana_data.csv"))
cocoa_clean <- cocoa2 %>%
  rename(Date = `Date`,
         Price = `ICCO.daily.price..US..tonne.`) %>%
  mutate(
```

```
    Date = dmy(Date),
    Price = as.numeric(gsub(",", "", Price))
  ) %>%
  arrange(Date)

cocoa_monthly <- cocoa_clean %>%
  mutate(Month = floor_date(Date, "month")) %>%
  group_by(Month) %>%
  summarise(Avg_Price = mean(Price, na.rm = TRUE)) %>%
  ungroup()

ghana_clean <- ghana2 %>%
  mutate(
    DATE = ymd(DATE),
    PRCP = replace_na(PRCP, 0),
    TMAX = na.locf(TMAX, na.rm = FALSE),
    TMIN = na.locf(TMIN, na.rm = FALSE),
    TAVG = as.numeric(TAVG)
  ) %>%
  filter(!is.na(DATE) & !is.na(TAVG))

ghana_daily <- ghana_clean %>%
  group_by(DATE) %>%
  summarise(
    TAVG = mean(TAVG, na.rm = TRUE),
    TMAX = mean(TMAX, na.rm = TRUE),
    TMIN = mean(TMIN, na.rm = TRUE),
    PRCP = sum(PRCP, na.rm = TRUE),
    .groups = 'drop'
  )

ghana_monthly <- ghana_daily %>%
  mutate(Month = floor_date(DATE, "month")) %>%
  group_by(Month) %>%
  summarise(
    Avg_TAVG = mean(TAVG, na.rm = TRUE),
    Avg_TMAX = mean(TMAX, na.rm = TRUE),
    Avg_TMIN = mean(TMIN, na.rm = TRUE),
    Total_PRCP = sum(PRCP, na.rm = TRUE),
    .groups = 'drop'
  )

combined_data <- inner_join(cocoa_monthly, ghana_monthly, by = "Month")
```

```
\#\#\# Merge and Clean Monthly Data
data <- combined_data %>%
  mutate(
    log_price = log(Avg_Price),
    diff_log_price = c(NA, diff(log_price))
  ) %>%
  drop_na()




\#\# ----------------------------------------------------------------------
\#| label: cocoa-price
\#| fig-cap: The Visualization of Monthly Cocoa Price
\#| echo: false
\#| message: false
\#| warning: false

library(gridExtra)
\#\#\# ------------------EDA and Time Series Decomposition----------------------
plot_price = ggplot(combined_data, aes(x = Month, y = Avg_Price)) +
  geom_line(color = "forestgreen") +
  labs(y = "Price", x = "Date") +
  theme_minimal()

plot_price_log = ggplot(data, aes(x = Month)) +
  geom_line(aes(y = log_price), color = "tomato") +
  labs(y = "Log(Price)", x = "Date") +
  theme_minimal()

plot_price_diff = ggplot(data, aes(x = Month)) +
  geom_line(aes(y = diff_log_price), color = "royalblue") +
  labs(y = "Difference Log(Price)", x = "Date") +
  theme_minimal()

grid.arrange(plot_price, plot_price_log, plot_price_diff, nrow=2)



\#\# ----------------------------------------------------------------------
\#| label: exploratory-data-analysis
\#| echo: false
\#| message: false
\#| warning: false
```

```
\#| fig-cap: STL Decomposition of Time Series Data (1993-2025)

ts_log_price <- ts(data$log_price, start = c(1994, 10), frequency = 12)
decomp <- stl(ts_log_price, s.window = "periodic")
plot_a = plot(decomp)
\# mean(data$Avg_Price)
plot_a


\#\# ------------------------------------------------------------------------
\#| label: ets-model
\#| echo: false
\#| message: false
\#| warning: false

\#\#\# Split Data into Training and Testing Sets
train_size <- floor(0.8 * nrow(data))
train <- data[1:train_size, ]
test <- data[(train_size + 1):nrow(data), ]
\#\# ETS Model
ets_auto_model <- ets(train$diff_log_price)  \# Auto ETS (default)
ets_explicit_model <- ets(train$diff_log_price, model = "ZZZ")  \# Explicit auto ETS

\# Forecast using ETS Models
ets_auto_forecast <- forecast(ets_auto_model, h = nrow(test))
ets_explicit_forecast <- forecast(ets_explicit_model, h = nrow(test))

\# Evaluate Forecast Accuracy
ets_auto_accuracy <- accuracy(ets_auto_forecast, test$diff_log_price)
ets_explicit_accuracy <- accuracy(ets_explicit_forecast, test$diff_log_price)


\#\# ------------------------------------------------------------------------
\#| label: arima-sarima-model
\#| echo: false
\#| message: false
\#| warning: false

\# ---- ARIMAX model (non-seasonal) ----
\# Prepare external regressors
train_xreg <- train %>%
  select(Total_PRCP, Avg_TAVG, Avg_TMAX, Avg_TMIN) %>%
  as.matrix()
```

```
test_xreg <- test %>%
  select(Total_PRCP, Avg_TAVG, Avg_TMAX, Avg_TMIN) %>%
  as.matrix()

arimax_model <- auto.arima(train$diff_log_price, xreg = train_xreg, seasonal = FALSE)

\# Forecast using ARIMAX
arimax_forecast <- forecast(arimax_model, xreg = test_xreg, h = nrow(test))

\# Evaluate ARIMAX model
arimax_accuracy <- accuracy(arimax_forecast, test$diff_log_price)

\#\# SARIMAX Model
\# Fit SARIMAX model
sarimax_model <- auto.arima(train$diff_log_price, xreg = train_xreg, seasonal = TRUE)

\# Forecast using SARIMAX model
sarimax_forecast <- forecast(sarimax_model, xreg = test_xreg, h = nrow(test))

\# Evaluate SARIMAX model accuracy
sarimax_accuracy <- accuracy(sarimax_forecast, test$diff_log_price)


\#\# ----model-output, message = FALSE, warning = FALSE--------------------------

\#\# Model Performace
cat("ETS Model 1 Performance:\n")
print(ets_auto_accuracy)

cat("ETS Model 2 Performance:\n")
print(ets_explicit_model)

cat("ARIMAX Model Performance:\n")
print(arimax_accuracy)

cat("SARIMAX Model Performance:\n")
print(sarimax_accuracy)


\#\# ----model-transform, echo = FALSE, message = FALSE, warning = FALSE----------
```

```
### Back-transform forecasted values

# Helper function to reconstruct log prices from differences
reconstruct_log_prices <- function(last_log, diffs) {
  cumsum(c(last_log, diffs))[-1]
}

# Get last observed log price from training set
last_log_price <- tail(train$log_price, 1)

# Optional: save forecast dates (if needed for plotting)
forecast_dates <- test$Month

# Reconstruct log-scale forecasts
ets_auto_log_forecast     <- reconstruct_log_prices(last_log_price, ets_auto_forecast$mean)
ets_explicit_log_forecast <- reconstruct_log_prices(last_log_price, ets_explicit_forecast$mea
arimax_log_forecast       <- reconstruct_log_prices(last_log_price, arimax_forecast$mean)
sarimax_log_forecast      <- reconstruct_log_prices(last_log_price, sarimax_forecast$mean)

# Convert log forecasts back to original price scale
ets_auto_price_forecast     <- exp(ets_auto_log_forecast)
ets_explicit_price_forecast <- exp(ets_explicit_log_forecast)
arimax_price_forecast       <- exp(arimax_log_forecast)
sarimax_price_forecast      <- exp(sarimax_log_forecast)

forecast_df <- bind_rows(
  tibble(Date = forecast_dates, Forecast = ets_auto_price_forecast, Model = "ETS Model 1"),
  tibble(Date = forecast_dates, Forecast = ets_explicit_price_forecast, Model = "ETS Model 2
  tibble(Date = forecast_dates, Forecast = arimax_price_forecast, Model = "ARIMAX"),
  tibble(Date = forecast_dates, Forecast = sarimax_price_forecast, Model = "SARIMAX")
) %>% drop_na()



## ----plot-forecast-vs-actual, echo = FALSE, message = FALSE, warning = FALSE----

data <- data %>% rename(Date = Month, Price = Avg_Price)

ggplot() +
  # Actual cocoa prices line
  geom_line(data = data, aes(x = Date, y = Price), color = "black", linewidth = 1.2) +

  # Forecast lines by model
```

```r
  geom_line(data = forecast_df, aes(x = Date, y = Forecast, color = Model, linetype = Model)

  \# Labels and theme
  labs(
    title = "Monthly Forecasts vs Actual Cocoa Prices",
    y = "Price",
    x = "Date"
  ) +
  theme_minimal() +
  theme(legend.position = "bottom") +

  \# Manual color and linetype mappings
  scale_color_manual(values = c(
    "ETS Model 1" = "blue",
    "ETS Model 2" = "green",
    "ARIMAX"      = "orange",
    "SARIMAX"     = "purple"
  )) +
  scale_linetype_manual(values = c(
    "ETS Model 1" = "solid",
    "ETS Model 2" = "dashed",
    "ARIMAX"      = "twodash",
    "SARIMAX"     = "dotdash"
  ))


\#\# ----random-forest, echo = FALSE, message = FALSE, warning = FALSE------------
\#----------------------------Random Forest ------------------------------

ml_data <- combined_data %>%
  mutate(
    Lag1 = lag(Avg_Price, 1),
    Lag2 = lag(Avg_Price, 2),
    Month_Num = month(Month),
    Year = year(Month)
  ) %>%
  drop_na()


split_index <- floor(0.8 * nrow(ml_data))
train_data <- ml_data[1:split_index, ]
test_data <- ml_data[(split_index+1):nrow(ml_data), ]
```

```r
rf_model <- randomForest(
  Avg_Price ~ Avg_TAVG + Total_PRCP + Lag1 + Lag2 + Month_Num,
  data = train_data,
  ntree = 500,
  importance = TRUE
)

pred_rf <- predict(rf_model, newdata = test_data)

results <- data.frame(
  Date = test_data$Month,
  Actual = test_data$Avg_Price,
  Predicted = pred_rf
)

rmse_rf <- rmse(pred_rf, test_data$Avg_Price)
mae_rf <- mae(pred_rf, test_data$Avg_Price)

cat("Random Forest RMSE:", round(rmse_rf, 2), "\n")
cat("Random Forest MAE :", round(mae_rf, 2), "\n")


\#\# ----fig-rf-full-fit, echo = FALSE, message = FALSE, warning = FALSE----------

full_results <- ml_data %>%
  mutate(
    Predicted = predict(rf_model, newdata = ml_data)
  ) %>%
  select(Month, Actual = Avg_Price, Predicted)

ggplot(full_results, aes(x = Month)) +
  geom_line(aes(y = Actual, color = "Actual")) +
  geom_line(aes(y = Predicted, color = "Predicted")) +
  scale_color_manual(values = c("Actual" = "black", "Predicted" = "blue")) +
  labs(title = "Random Forest Prediction vs Actual (Monthly, Full Period)",
       x = "Month", y = "Cocoa Price (US$/tonne)",
       color = "Legend") +
  theme_minimal() +
  theme(legend.position = c(0.05, 0.95),
        legend.justification = c("left", "top"))


\#\# ----xgboost, echo = FALSE, message = FALSE, warning = FALSE-----------------
```

```
\#| fig-cap: Rolling XGBoost Forecast vs Actual Cocoa Prices
\#| label: fig-xgboost

\#\#\# --- Load Libraries ---
library(xgboost)
library(dplyr)
library(lubridate)
library(tidyr)
library(ggplot2)

\#\#\# --- Create Lag Features Function ---
generate_lags <- function(data, lags = c(1:12, 24)) {
  for (lag in lags) {
    data[[paste0("lag_", lag)]] <- dplyr::lag(data$log_price, lag)
  }
  return(data)
}

\#\#\# --- Load and Prepare Data ---
\# Assume `data` contains: Date, Price, PRCP, TAVG, TMAX, TMIN
data <- data %>%
  arrange(Date) %>%
  mutate(
    log_price = log(Price),
    month = month(Date),
    year = year(Date),
    time_index = 1:n()
  ) %>%
  generate_lags() %>%
  drop_na()

\#\#\# --- Setup Result Storage ---
forecast_start <- 120  \# start forecasting after this many rows (10 years of monthly data)
forecast_end <- nrow(data)
results <- data.frame(
  Date = as.Date(character()),
  Actual = numeric(),
  Predicted = numeric()
)

\#\#\# --- Rolling Forecast Loop ---
for (i in forecast_start:(forecast_end - 1)) {
  train_data <- data[1:i, ]
```

```r
  test_data <- data[i + 1, , drop = FALSE]

  \# Skip if missing
  if (nrow(test_data) == 0 || any(is.na(test_data))) next

  \# Training matrix
  x_train <- train_data %>%
    select(starts_with("lag_"), Total_PRCP, Avg_TAVG, Avg_TMAX, Avg_TMIN, month, year, time_
  y_train <- train_data$log_price

  dtrain <- xgb.DMatrix(data = as.matrix(x_train), label = y_train)

  \# Fit XGBoost model
  xgb_model <- xgboost(
    data = dtrain,
    nrounds = 300,
    eta = 0.05,
    max_depth = 10,
    subsample = 0.8,
    colsample_bytree = 0.8,
    objective = "reg:squarederror",
    verbose = 0
  )

  \# Predict next step
  x_test <- test_data %>%
    select(starts_with("lag_"), Total_PRCP, Avg_TAVG, Avg_TMAX, Avg_TMIN, month, year, time_
  dtest <- xgb.DMatrix(data = as.matrix(x_test))
  pred_log <- predict(xgb_model, dtest)
  pred_price <- exp(pred_log)

  \# Save results
  results <- rbind(results, data.frame(
    Date = test_data$Date,
    Actual = exp(test_data$log_price),
    Predicted = pred_price
  ))
}

\#\#\# --- Prepare Data for Plotting ---
results_long <- results %>%
  pivot_longer(cols = c("Actual", "Predicted"), names_to = "Type", values_to = "Price")
```

```
\#\#\# --- Plot with Legend ---
ggplot(results_long, aes(x = Date, y = Price, color = Type)) +
  geom_line(linewidth = 1.2) +
  scale_color_manual(values = c("Actual" = "red", "Predicted" = "blue")) +
  labs(
    title = "Rolling XGBoost Forecast vs Actual Cocoa Prices",
    y = "Price", x = "Date",
    color = "Legend"
  ) +
  theme_minimal()



\#\# ---------------------------------------------------------------------
\#\#\# --- Evaluation ---
rmse <- sqrt(mean((results$Actual - results$Predicted)^2))
cat("Rolling XGBoost RMSE:", round(rmse, 2), "\n")
mape <- mean(abs((results$Actual - results$Predicted) / results$Actual)) * 100
cat("MAPE:", round(mape, 2), "%\n")



\#\# ---------------------------------------------------------------------

\# error table



\#\# ----message=FALSE, warning=FALSE, echo=FALSE, results='hide'----------------
\# Save the code file silently without showing output
\# dummy <- knitr::purl("STA457_Project.qmd", output = "appendix_code.R")



\#\# ----results='asis', echo=FALSE, message=FALSE, warning=FALSE----------------
\#| echo: false
\#| eval: true
\# system("cat appendix_code.R")
```